Tejas AI — AI Factory Elevated Pitch

Test autonomous control in the twin · Repair in VR · Deploy to real on your terms

> **Audience:** AI factory operators, GPU colo builders, hyperscaler-adjacent edge DCs, NVIDIA-class training halls.
> **Angle:** Phaidra proved AI factories need autonomous cooling — but nobody lets you **practice** before you touch live hardware. Tejas is the **testing environment**. VR repair cuts downtime and cost.
> **Demo:** `cd tejas-twin && ./run.sh` → `/datacenter` + `/vr?enter=1`

The one line (new)

> **Tejas gives every AI factory a full digital twin testing ground — train autonomous cooling, stress heat waves, and repair faults in VR — so your team graduates to the real hall with proof, not prayer.**

Why this pitch is different

Old pitch	**AI Factory pitch**
“We save 20% cooling ₹”	“Your AI factory **requires** autonomous control — we let you **test it safely first**”
Chennai heat-wave story	**GPU hall at 47°C** — same physics as H100/B200 training load
AR fix as side feature	**VR repair is the hero** — less downtime, less cost, technician never lost
We deploy for you	**You test in twin → shadow → you own the real env** when ready

The elevated narrative (3 beats)

```
1. AI FACTORIES CHANGED THE RULES
   Rack power is a leading indicator. Autonomous control in <10s.
   Phaidra runs NVIDIA halls. Everyone else is still on 1990s BMS
logic.

2. YOU CANNOT LEARN ON LIVE GPUs
   One wrong setpoint = millions in failed jobs.
   Tejas = the testing environment: twin → shadow → advise → control.
   Break chillers, heat waves, GPU-16 faults — in simulation. Zero
risk.
```

```
3. REPAIR IN VR = LESS DOWNTIME, LESS COST
   Don't fly a senior tech across the country for a filter.
   Walk the hall in VR, inspect the rack, see the exploded design,
   fix in AR — 25 minutes, not 25 hours of outage.
```

SLIDE DECK — 14 slides (copy to PowerPoint or use `Tejas-AI-Factory-Deck.pptx`)

Slide 1 — TITLE

Visual: Dark navy. Single glowing rack line receding into VR perspective.

Headline: # TEJAS AI ### The AI factory testing environment

Sub: Train autonomous control · Stress the twin · Repair in VR · Deploy to real when you're ready

Footer: तेजस् · On-prem · Sovereign · Built for GPU halls

Speaker note: Don't open with savings. Open with fear: *you cannot learn autonomous cooling on live H100s.*

Slide 2 — THE AI FACTORY PROBLEM

Headline: AI factories broke the old playbook

Bullets: - Rack power spikes **before** inlet temp moves — you have **seconds**, not minutes - Training jobs run **weeks** — one cooling trip = **₹2+ Cr** in lost compute + SLA hell - Phaidra proved autonomous control works — on **NVIDIA-scale** AI factories - Everyone else still runs **rule-based BMS** set in the 1990s

Stat box: `30–40%` of hall power = cooling · `47°C` ambient = margin gone

Speaker note: "The giants proved it works. The mid-market has no way to **practice** before they flip the switch."

Slide 3 — YOU CANNOT TEST ON LIVE SILICON

Headline: The real hall is not a lab

Three columns:

Test on live hall	Test on Tejas twin
❌ Failed job = customer churn	✅ Heat wave in one drag
❌ Wrong CHW setpoint = chip throttle	✅ Chiller trip — recover in sim
❌ No rollback story for the board	✅ Shadow mode: log would-do vs BMS
❌ Insurance / MoC nightmare	✅ Trust ladder with hard revert

Pull quote: > *"Would you let an untested AI drive your GPUs on Thursday's 47°C day? Neither would we."*

Slide 4 — TEJAS = THE TESTING ENVIRONMENT

Headline: A full AI factory — on one laptop — before you touch hardware

Diagram:

```
YOUR REAL HALL          TEJAS TESTING ENV           GRADUATION PATH
──────────────         ─────────────────           ───────────────
BMS (read-only)   →    Digital twin (physics)  →    You run shadow
Historian export  →    Autonomous brain        →    You approve advise
GPU telemetry     →    47°C stress tests       →    You take control
                       VR walk + repair             when YOUR data
proves it
```

What they get day one: - Calibrated twin (topology + historian — not a drawing toy) - AI vs baseline running **live** — PUE 1.3 vs 1.5 at 47°C - **Zero write access** until they choose to graduate

Speaker note: "We don't ask you to trust us. We give you a **sandbox** that looks like your hall."

Slide 5 — LIVE DEMO BREAK

Visual: Full-bleed screenshot placeholder. Big text: **DEMO NOW — 2 minutes**

Demo script (in order): 1. `/datacenter` — Autonomous Control card · PUE live 2. Weather → **47°C** — baseline panics, Tejas holds 3. *"Is any machine going down?"* → **GPU-16** pulses red 4. Open **VR QR** on phone → `/vr?enter=1`

Say: "This is the testing environment. Same physics. Your GPUs are not at risk."

Slide 6 — AUTONOMOUS CONTROL (THE REQUIREMENT)

Headline: What the testing environment trains

Bullets: - **Feed-forward control** — acts on rack power + weather **before** inlet spikes - Grid-search today → MPC/RL in production — **same trust ladder** - Explains every move: *“Heat wave: running supply warmer to protect COP”* - Dual run: **Tejas vs dumb BMS** — savings gap visible in real time

Comparison table:

	Baseline BMS	Tejas in twin
47°C response	Reactive · power spike	Predictive · holds band
PUE	~1.50	~1.30
Operator	Dashboard training	*“Push outside to 47”*

Future in real hall: Shadow 30 days → they see **their** numbers → they flip control themselves.

Slide 7 — THE TRUST LADDER (YOU GRADUATE, NOT US)

Headline: From testing env to real hall — on your schedule

```
TWIN           SHADOW           ADVISORY         SUPERVISED
AUTONOMOUS
(test here)    (read-only)      (you apply)      (you confirm)     (you
own it)
     |              |                 |                 |
|
  Break things   Your historian   Your operator    Your MoC
Your control
  in sim         proves ₹         clicks apply     sign-off
policy
```

Key line: > **We hand you the keys when your data says yes — not when our sales deck says yes.**

Commercial: Shadow pilot = low fee or free · Savings-share only after **they** choose advisory+

Slide 8 — HERO: REPAIR IN VR

Headline: Less downtime. Less cost. The technician never walks blind.

Visual: Split — left: empty aisle + flight ticket ₹₹₹ / right: phone VR + GPU-16 reticle

The cost of the old way: | Old repair path | Tejas VR + AR path | |—|—| | Alarm at 2:14 AM | GPU-16 flagged **14 months early** | | Senior tech on a plane | **VR walk** to exact rack — any junior on site | | Open the rack blind | **Exploded design** — filter, fan tray, busbar | | 4-hour diagnosis | **25-minute** fix with parts list pre-staged | | ₹2+ Cr outage | **Planned window** — hall never tripped |

Stat: `14 months` early detection · `25 min` on-site fix · `₹8.3L/yr` waste eliminated per rack-class fault

Speaker note: THIS is the emotional peak. Put the phone in their hand. Let them **stand inside** the hall.

Slide 9 — VR REPAIR FLOW (SHOW THIS LIVE)

Headline: Walk in · Pick · Inspect · Fix · Blue

Steps (match product): 1. **Scan QR** → stand inside the live twin (`/vr`) 2. **Joystick walk** to Row 2 — GPU-16 pulsing red 3. **Reticle pick** → SERVICE NEEDED · inlet · load · CFM 4. **Aspect lenses** → Cooling: CRAC-2 filter + fan tray highlighted 5. **Field AR** (`/ar`) → 4-step checklist on the **real** machine 6. Rack turns **blue** — work order closed — desktop syncs

Say: > *"Your tech in Chennai doesn't need a senior from Frankfurt. They **walked the fault in VR** last Tuesday."*

Downtime math (say aloud): - Unplanned GPU hall trip: ₹**2+ Cr** - VR-guided planned fix: **25 min** · parts known · **zero** surprise

---

Slide 10 — PREDICT BEFORE THE TRIP

Headline: GPU-16 — the story every AI factory has, but doesn't see

Bullets: - Inlet **3.4°C hotter** than load explains — 14-month creep - Root cause: CRAC-2 filter + rear fan-tray bearing - Tejas drafts the **work order + email** with exact parts - Caught in the **testing twin** first — then validated on real historian in shadow

Chat line: *"Is any machine going down?"*

VR line: Reticle on GPU-16 → *"SERVICE NEEDED"*

Link to repair: Prediction in twin → practice fix in VR → execute in AR on site

---

Slide 11 — NUMBERS (AI FACTORY SCALE)

Headline: What the testing environment proves — before one setpoint hits silicon

12 MW GPU hall · India tariff:

Line	Annual
Cooling energy saved (20–28%)	₹**1.8–3.2 Cr**
One avoided outage	₹**2+ Cr**
One VR-guided fix vs emergency	₹**8.3L+** waste cut

Line	Annual
Fleet of 50 AI halls	**₹90–160 Cr**

Water / WUE: Chennai-class — licence to operate, not just efficiency

Say: “Shadow mode puts **your** historian on this slide — not our simulation.”

Slide 12 — WHO THIS IS FOR

Headline: Built for the halls Phaidra doesn’t serve — yet

Buyer	Why Tejas testing env
GPU colo / edge AI DC	Can’t afford a Phaidra pilot · need sovereign on-prem
Enterprise AI lab	Internal team must **own** control policy — we train them in twin
DC operator adding GPU block	Liquid + air mixed — test before go-live
ESCO / integrator	White-label testing sandbox for customers

Not for: Greenfield with no telemetry yet (advisory-only twin first)

Slide 13 — COMMERCIAL

Headline: Pay for proof — not promises

Phase	What they buy	Price logic
Testing env / twin build	Calibrated sandbox + VR + AR	Paid on-ramp (credited to Y1)
Shadow (30 days)	Read-only · would-do report	Low / free — proves ₹
Advisory + control	They graduate when ready	**Savings-share 20–40%**

Land-and-expand: One hall testing env → campus → fleet panel (50 sites)

Slide 14 — CLOSE

Visual: Dark. VR silhouette in aisle. Blue corrected rack glow.

Headline: # Don’t learn autonomous cooling on live GPUs. # Learn it in the twin. Repair it in VR. Own the real hall when you’re ready.

Three lines: 1. **Test** autonomous control — heat waves, load spikes, chiller trips — zero risk 2. **Repair** in VR — less downtime, less cost, parts before the plane ticket 3. **Graduate** on your data — shadow → advise → control — you flip the switch

CTA: > *“Give us 30 days read-only on your historian. We’ll show you the testing environment built from your hall — and one GPU rack you’d have lost on a 47°C Thursday.”*

Contact: [email] · [calendar]

60-SECOND ELEVATED SCRIPT (AI Factory + VR)

HOOK (0:00–0:18)

> Every AI factory now needs **autonomous cooling** — rack power moves in seconds, not minutes. Phaidra proved it on NVIDIA halls. But **you cannot practice that on live GPUs** — one wrong setpoint kills a three-week training job.

WOW (0:18–0:42)

> *Open `/datacenter` then hand them VR.* Tejas is the **testing environment** — a full physics twin on one laptop. Watch: **47 degrees** — baseline panics, our brain holds PUE 1.3. *“Is any machine going down?”* — GPU-16, fourteen months early. Now **stand inside** — scan this QR — you’re in the hall. Walk to the red rack. **That’s the repair before the outage** — exploded design, exact parts, twenty-five minutes, not two fourteen AM.

CLOSE (0:42–1:00)

> Test autonomous control here. Graduate to your real hall when **your** shadow data says yes. Repair in VR — **less downtime, less cost**. We don’t sell dashboards. We sell **the sandbox where your team earns the right to flip the switch.**

3-MINUTE MEETING FLOW

Min	Do	Say
0:00	Slide 2–3	AI factory problem · can’t test on live silicon
0:45	Slide 4	Tejas = testing environment
1:00	**Demo** datacenter 47°C + GPU-16	Autonomous control in twin
1:45	**Demo VR** — hand phone, walk to GPU-16	Repair hero moment
2:15	Slide 8–10	Downtime cost · VR flow · predict

Min	Do	Say
2:45	Slide 11 + 13	Numbers · shadow pilot CTA
3:00	Slide 14	Close · calendar

DEMO CHECKLIST (AI Factory pitch)

- ☐ `run.sh` same day · weather pre-loaded 43°C
- ☐ GPU-16 fault visible before VR handoff
- ☐ **VR QR ready** on laptop screen (`/vr` link from datacenter)
- ☐ Phone on same Wi-Fi · HTTPS cert accepted
- ☐ Optional: AR QR for fix checklist after VR pick
- ☐ Shadow pilot one-pager printed with **their** hall name
- ☐ Backup screen recording if room Wi-Fi fails

OBJECTION HANDLERS (AI Factory specific)

Objection	Response
"We already have Phaidra / Schneider"	"They run production. We give your **internal team a testing sandbox** + VR repair — sovereign, on-prem, before you depend on anyone."
"We won't let AI near GPUs"	"Neither would we — **day one is read-only shadow**. The testing env breaks things in **simulation**, not silicon."
"VR is a gimmick"	"A single unplanned trip is ₹**2 Cr**. VR is how your Chennai tech **practices** the fix before Frankfurt gets on a plane."
"Our team will build this"	"Build the twin infra + safety supervisor + VR repair loop — or **start testing Monday** on one laptop and graduate on your historian in 30 days."

Deck file: `Tejas-AI-Factory-Deck.pptx`
Pair with: `60-second-elevated-pitch.md` (general) · `pitch.md` (full playbook)