

Tejas AI — Complete Pitch Kit

Everything you need to sell, demo, fundraise, and close.

Companion docs: `story.md` (narrative) · `project.md` (strategy) · `demo.md` (two-screen demo) · **Slide deck:** `Tejas-AI-Pitch-Deck.pptx`

0. Before you open your mouth

Asset	Where	When to use
Slide deck	<code>pitch/Tejas-AI-Pitch-Deck.pptx</code>	Investor meeting, board intro, conference
One-pager	<code>pitch/one-pager.md</code>	Email attachment, WhatsApp, leave-behind
Live demo	<code>cd tejas-twin && ./run.sh → https://localhost:7878</code>	Every serious meeting — show, don't tell
Story	<code>story.md</code>	Deep-dive with technical buyers, case-study framing
Recorded backup	Record a perfect 3-min screen capture	Room fights you, bad Wi-Fi, async follow-up

Golden rule: Lead with the **live demo** or a **recorded clip**. The deck supports the demo; it does not replace it.

1. One line (memorise this)

Tejas AI is a self-learning agent that runs a facility's cooling at peak efficiency — cutting energy 20–40% — trained safely on a digital twin, deployed without ripping out hardware, and operated by talking to it.

Name: Tejas (तेजस्) — radiance, sharpness. India's indigenous fighter jet. Built to tame heat.

2. Elevator pitches

★ **Product definition (every feature):** `./PRODUCT.md`

★ **Best 60-second pitch (30·40·30 + live wow):** `60-second-elevated-pitch.md`

10 seconds

“We put an AI brain on top of your existing cooling system. It predicts heat waves, cuts cooling energy 20–40%, and your operator just talks to it — no dashboard training.”

30 seconds (investor)

“Cooling is 30–40% of a data center’s power bill, and almost all of it runs on 1990s rule-based logic. DeepMind and Phaidra proved AI can do better — but they only serve hyperscalers. Tejas is the on-prem, sovereign cooling brain for everyone else: cold chain, mid-size DCs, factories. We prove savings on a digital twin first, then shadow your real plant read-only for 30 days. Zero capex objection — we take a share of verified savings.”

30 seconds (customer / facility owner)

“Your BMS reacts after heat builds. Tejas predicts it — weather, load, thermal drift — and optimises every minute inside your safety band. We connect read-only to what you already have. No rip-and-replace. Month one: shadow mode shows what we would have saved on your actual data. Month two: you decide if you want us to advise or run it. You only pay from money you weren’t going to have.”

2 minutes (full hook — use after demo)

“Last summer, India’s longest heat wave on record had engineers warning of grid blackouts as AC load spiked. In Chennai’s data-center corridor, cooling towers were at the edge of their envelope — same physics as a hall trip: ambient up, chiller COP down, margin gone.

The cruel part: every predictive signal was already in the building. BMS trends. Historian data. Weather forecasts. Nobody was reading it as a story.

Tejas reads that story. We build a calibrated digital twin of your facility — racks, chillers, cold rooms, separators — tune it against your historian until simulation matches reality. Then we train the AI in simulation, deploy in shadow mode, and climb a trust ladder: watch → advise → control. The operator asks in English, Hindi, or Tamil: *‘What happens if Thursday hits 47 degrees?’* — and gets numbers, not shrugs.

On a 47°C day we hold the same safety band at PUE ~1.30 vs baseline ~1.50. GPU-16-style predictive maintenance flags a degrading rack fourteen months before inspection. Field AR guides the technician to the exact part.

Everything you just saw runs on one laptop, offline. The twin is the method — not a fake. That’s how we sell to conservative operators who will never let an unproven AI near critical cooling without proof first.”

3. The opening question (use on slide 1 or as your first words)

“What did your building know last night that you didn’t?”

Pause. Let it land. Then:

“Your BMS knew the chiller was short-cycling. Your historian knew inlet temps were drifting. Your weather feed knew Thursday would be 47°C. The data was there. **The story wasn’t.**”

4. Investor pitch — 12-minute structure

Min	Section	You say	Show
0:00	Hook	Opening question + Chennai heat-wave context	Slide 2–3 or jump straight to demo
1:00	Problem	30–40% of DC power is cooling; dumb BMS; India 40–48°C; Phaidra ignores mid-market	Problem slides
2:30	Live demo	90-second datacenter tour (see §8)	/datacenter — weather to 47°C, GPU-16, fleet
4:00	Solution	Predict + optimise + explain + act; twin-first; trust ladder	Trust ladder slide
5:00	Product	Platform: DC, factory, geothermal, cold chain, AR, VR, ingestion	Admin panel quick tour
6:30	Traction / proof	Simulation validated; real anchors (Chennai, Geysers, Renesas); pilot-ready ingestion wizard	story.md anchors
7:30	Market	India cold chain + DC boom + buildings; global extension	Market slide
8:30	Business model	Savings-share primary; SaaS; twin-build engagement	Pricing slide
9:30	Moat	Calibrated twin + safe control policy + per-site operational data; on-prem sovereign	Competition slide
10:30	GTM	Demo wedge → shadow pilot → savings-share → partnerships (BEE, ESCOs, cold-chain assoc.)	GTM slide

Min	Section	You say	Show
11:30	The ask	[Fill in: round size, use of funds, milestones]	Ask slide
12:00	Close	“We sell rupees off the power bill — for everyone the giants ignore.”	Title slide

Honest framing (say it — builds trust): > “The demo is validated in simulation on real Indian weather. The pilot proves it on your facility: read-only shadow for 30 days, savings on your data, before we ever touch a setpoint.”

5. Customer pitch — 20-minute structure

Min	Section	Goal
0:00	Their pain	Ask: “What was your worst cooling incident / power bill surprise last summer?” Listen.
3:00	Mirror	Reflect their answer with Chennai / cold-chain / fab story from <code>story.md</code>
5:00	Demo	Full 3-minute script (§8) tailored to their vertical
8:00	How we’d deploy here	Shadow → advisory → control; BACnet/Modbus/Redfish; no rip-and-replace
12:00	ROI	Their numbers: MW, tariff, cooling %, downtime cost → §9 calculator
15:00	Pilot proposal	30-day shadow, deliverables, success criteria (§11)
18:00	Pricing	Savings-share: “You pay only from money you weren’t going to have”
20:00	Next step	Calendar hold for site walk + historian export checklist

Discovery questions (write answers on the call): 1. Facility type? MW / tons cooling? Air vs liquid? 2. BMS vendor? Historian? Last calibration? 3. Biggest pain: energy bill, downtime,

water, staffing, SLA? 4. Who signs? Who operates daily? 5. Any incident in last 24 months? Cost? 6. Current PUE or refrigeration kWh/ton if known?

6. Real-world anchors (credibility — with links)

Use these to ground the story. Demo characters (GPU-16, SEP-2, PRESS-7) are **composites**, not claims about one specific incident.

Vertical	Real anchor	Link
Data center	Chennai OMR/Siruseri corridor; CtrlS Tier-IV	CtrlS Chennai
Heat stress	2024 Indian heat wave — longest on record	Wikipedia · Guardian
Geothermal	The Geysers, CA — reinjection after steam depletion	Wikipedia
Factory	Renesas Naka fab fire, Mar 2021	Chip shortage
Cold chain	₹92,651 crore annual post-harvest loss	IBEF · ICRIER

7. Competitive positioning

vs	Their weakness	Tejas answer
Status quo BMS	Reacts; static rules; no story	Predicts; learns; explains every action
Phaidra / foreign AI	Hyperscaler-only; cloud; premium	Mid-market; on-prem; savings-share; India-tuned
Schneider / Vertiv	Hardware + dashboards	AI control layer on top — no rip-and-replace
Build in-house	No twin infra, no ML ops	Off-the-shelf brain; shadow proof in 30 days
ESCO / audit-only	Report, no closed loop	Twin + autonomous control + AR field execution

Moat (what they cannot copy in a quarter): 1. Calibrated digital twin per facility (topology, not asset count) 2. Safe control policy + Safety Supervisor architecture 3. Proprietary operational data accumulated per site 4. Vernacular voice + on-prem sovereign stack

8. Demo scripts

8A. Flagship — 90 seconds (investor / judge)

Prep: cd tejas-twin && ./run.sh → https://localhost:7878 → Admin → **Data Center**

Step	Action	Say
1	Point at Autonomous Control card	“It’s already running. Policy version, confidence, live PUE 1.3 vs 1.5 baseline.”
2	Drag Weather → 47°C (or type in chat)	“Heat wave. Sky goes red. Baseline panics. Tejas holds the band and widens savings.”
3	Chat: “ <i>How much electricity are we using?</i> ”	“Answers from live state — not a script.”
4	Chat: “ <i>Is any machine going down?</i> ”	“GPU-16 — degrading 14 months. Root cause, parts list, drafted service email. Rack pulses red in 3D.”
5	Chat: “ <i>Project savings across 50 sites.</i> ”	“Fleet panel — boardroom slide that signs the contract.”
6	Close	“One laptop. Offline. Twin is how we prove before we touch your hardware.”

8B. Full — 3 minutes (customer)

#	Say (voice or chat)	What happens
1	“ <i>Show me Chennai data center, AI versus baseline.</i> ”	Comparison view; sim running
2	“ <i>How much are we saving right now?</i> ”	Savings panel: ₹/day, kWh, CO ₂
3	“ <i>Heat wave — push outside to 47 degrees.</i> ”	Baseline power climbs; AI adapts; gap widens
4	“ <i>Are we still safe?</i> ”	Racks in band; “same safety, less power”
5	“ <i>Is any machine going down?</i> ”	GPU-16 story + service email
6	“ <i>Project across 50 sites.</i> ”	Fleet projection
7	(Hindi) “ <i>आज सबसे गरम यूनिट कौन सी है?</i> ”	Vernacular works
8	Open Field AR on phone (optional)	QR → navigate to GPU-16 → fix steps

#	Say (voice or chat)	What happens
9	Close	Shadow-mode pilot on their data — zero risk

8C. Vertical shortcuts

Vertical	Route	Hero moment
Brownfield DC	/ingest	6-step wizard: connect → discover → calibrate → shadow
Factory	/factory or /scan	Phone scan → twin in 10 min; PRESS-7 fault
Geothermal	/vr?type=geothermal	Walk to SEP-2; reinjection story
Cold chain	Weather + savings on factory preset	2–8°C band hold at 47°C ambient
VR wow	/vr	First-person hall walk; reticle-pick GPU-16

8D. Failure-proofing

- Voice fails → **type the same commands** in chat (built-in parser always works)
- Pre-warm sim before the meeting
- Rehearse exact phrases in §8B
- Keep recorded backup clip ready
- TEJAS_HTTP=1 only if HTTPS cert blocks you (camera AR needs HTTPS)

9. ROI talk track (fill on the call)

Single 12 MW hall, India industrial tariff (~₹8–12/kWh blended):

Line item	Conservative	Aggressive
Cooling energy saved (20–28%)	₹1.8 crore/yr	₹3.2 crore/yr
One avoided outage	₹2+ crore	₹5+ crore
One predictive fix (GPU-16 class)	₹8.3 lakh/yr waste eliminated	—
Fleet of 50 halls	₹90 crore/yr	₹160 crore/yr

Cold storage (₹65 lakh/month refrigeration bill): - 15–25% optimisation → ₹12–20 lakh/month → ₹1.4–2.4 crore/year - Spoilage reduction is upside on top (national scale: [IBEF](#))

Water (Chennai-class): Lower WUE = licence to operate, not just efficiency.

Carbon: ~0.71 kg CO₂/kWh India grid → hundreds of tonnes/hall/year — auditable for ESG.

Say: “We’ll give you a number on *your* historian in shadow mode — these are the class of outcomes we see in simulation.”

10. Trust ladder (deployment — draw this live)

PHASE 0 PHASE 4	PHASE 1	PHASE 2	PHASE 3
Build twin Autonomous Train in sim. AI runs Zero hardware. closed-loop. Hard limits + instant revert.	Shadow Read-only. Log would-do vs reality.	Advisory AI recommends. Human applies. Build trust.	Supervised AI writes. Human confirms. MoC sign-off.

Phase	Customer risk	Deliverable
Twin	None	Physics model — demoable today
Shadow	None	“What we would have saved — last 30 days”
Advisory	Low	Optimal setpoints on screen
Supervised	Medium	AI writes; human approves each change
Autonomous	Managed	Closed-loop inside envelope; revert to baseline on anomaly

11. Pilot proposal template (copy-paste to PDF/email)

Tejas AI — 30-Day Shadow Pilot

Scope: [Facility name], [city] — [DC / cold chain / plant type]

Week 1 — Connect (read-only) - BACnet / Modbus / OPC-UA / Redfish / SNMP / historian export - No write access to BMS - Asset discovery + topology map

Week 2–3 — Calibrate twin - Tune simulation vs measured temps (target $\pm 1^\circ\text{C}$) and energy ($\pm 8\%$) - Assumptions ledger + confidence gate

Week 4 — Shadow report - AI vs baseline: kWh, ₹, PUE, peak demand - Identified anomalies (e.g. airflow drift, short-cycling) - Recommended actions (not executed without approval)

Success criteria: - [] Credible savings range on customer's actual tariff - [] At least one predictive maintenance flag validated by facilities team - [] Customer sign-off to proceed to advisory phase

Commercials: Pilot fee ₹[X] credited toward first year **OR** pure savings-share after advisory phase — customer choice.

Data: On-prem only. No cloud egress. NDA standard.

12. Objection handling

Objection	Response
“We already have a BMS.”	“Perfect — we sit on top. We send better setpoints to hardware you already own. Month one is read-only.”
“AI can’t touch critical cooling.”	“Agreed. That’s why we shadow first, then advise, then control — with a Safety Supervisor and instant revert to your proven baseline.”
“Show me real savings, not a demo.”	“The demo proves the logic. The pilot proves your numbers. 30 days read-only on your historian — zero risk.”
“Phaidra already does this.”	“For NVIDIA-scale AI factories. Who optimises your cold stores and mid-size colos? We’re on-prem, India-tuned, savings-share.”
“Too expensive.”	“Savings-share: you pay only from verified savings. If we don’t save, you don’t pay.”
“Our team won’t use another dashboard.”	“They talk to it. Hindi, Tamil, English. ‘How much did we save today?’ — no training.”
“What if the AI breaks something?”	“Rate limits, hard safety envelope, sensor dropout → instant baseline. Every action logged. Facility is never worse off than before Tejas.”
“We need integration with [vendor].”	“BACnet, Modbus, OPC-UA, Redfish, SNMP — we ingest in the wizard. Show us your point list.”
“Data sovereignty?”	“On-prem by default. Ollama local LLM. Data never leaves the site.”

13. Email templates

13A. Cold outreach — data center operator

Subject: Chennai heat wave + your chiller short-cycle data

Hi [Name],

India's 2024 heat wave was the longest on record — and facilities in the Chennai corridor are running cooling at the edge of their envelope ([context](#)).

Most halls already have the signals in the BMS: short-cycling chillers, drifting inlet temps, rising WUE. Nobody reads it as a story.

Tejas AI is an on-prem cooling brain that: - Predicts heat waves and load spikes on a calibrated digital twin - Cuts cooling energy **20–28%** inside your existing safety band - Starts **read-only shadow mode** — no rip-and-replace

I'd like 20 minutes to show you a live demo on one laptop — then discuss a 30-day shadow pilot on your historian.

[Calendar link]

Best,
[Name]

13B. Cold outreach — cold chain

Subject: ₹65L/month refrigeration — what's your 47°C day plan?

Hi [Name],

Refrigeration is often 45% of a cold-storage electricity bill. India's post-harvest loss from missing cold infrastructure is estimated at **₹92,651 crore/year** — the room that drifts to 9°C for six hours loses the container.

Tejas AI holds your product band at the **warmest safe point** — not the coldest nervous point — and proves savings in shadow mode before touching setpoints.

Worth a 15-minute demo?

[Calendar link]

13C. Post-demo follow-up

Subject: Tejas demo + shadow pilot next steps

Hi [Name],

Thanks for your time today. As discussed:

1. **What you saw:** autonomous cooling at 47°C, PUE 1.3 vs 1.5, GPU-16 predictive maintenance, fleet projection

2. **What we'd do on your site:** 30-day read-only shadow on [BMS/historian] — savings report on your tariff
3. **Commercials:** savings-share option — pay only from verified savings

Attached: one-pager + pilot scope.

Can we schedule a 45-minute technical call with your facilities lead and historian expert?

Best,
[Name]

13D. Investor follow-up

Subject: Tejas AI — demo recording + deck

Hi [Name],

Following our conversation: Tejas is the on-prem cooling brain for the mid-market Phaidra ignores — cold chain, edge DCs, factories — proven twin-first, deployed shadow → control.

Attached: pitch deck, one-pager

Demo: [link to recording] or live at your office (one laptop, offline)

Round: [Raising ₹X pre-seed/seed] for [pilots / twin engine / GTM]

Happy to dive deeper on the calibrated-twin moat and savings-share unit economics.

Best,
[Name]

14. Pricing & packaging

Model	Who	Pitch line
Savings-share (20–40%)	Brownfield DC, cold chain	“Pay only from money you weren’t going to have”
SaaS per facility	Multi-site chains post-trust	Predictable recurring
Twin-build engagement	First site in a new logo	Paid on-ramp; funds calibration; produces ROI case
Pilot fee → credit	Risk-averse enterprise	De-risks procurement; credits to Y1

Land-and-expand: One hall → whole campus → fleet panel in product is the commercial vision.

15. Go-to-market sequence

1. **Clip the demo** — 30–60 sec heat-wave moment; post LinkedIn / X / YouTube Shorts
 2. **Target list** — Chennai DC operators, cold-chain assoc., hospital estates, food logistics
 3. **Shadow pilots** — 3–5 logos; publish anonymised case studies
 4. **Partnerships** — BEE, ESCOs, CtrlS-class operators, integrators
 5. **Content** — “What your BMS knew last night” blog series; open benchmark methodology
-

16. The ask (investor slide — customise)

Raising: ₹[X] [pre-seed / seed]

Use of funds: - 40% — Pilot deployments + calibration engineering - 30% — Twin engine + control policy hardening - 20% — GTM (demo tour, content, partnerships) - 10% — Legal, compliance, IP

12-month milestones: - [] 5 shadow pilots completed with documented savings ranges - [] 2 paid advisory/control contracts - [] Ingestion wizard production-hardened for BACnet + Redfish - [] ₹[X] ARR or savings-share revenue

17. Closing lines (pick one)

Customer:

> “Let’s shadow your facility for 30 days. Read-only. Your data. Your numbers. Then you decide.”

Investor:

> “The giants proved AI-run cooling works. We’re bringing it to everyone they ignore — on-prem, in India, starting with a demo you can run on this laptop tonight.”

Universal:

> “We don’t sell a dashboard. We sell rupees off the power bill — and we prove it before we touch a setpoint.”

18. Checklist — day of the meeting

- ☐ Laptop charged; `tejas-twin/run.sh` tested same day
- ☐ HTTPS cert accepted (or LAN IP for phone AR)
- ☐ OpenAI key or Ollama running (optional; built-in parser as fallback)
- ☐ Deck on USB + cloud
- ☐ One-pager printed or PDF ready
- ☐ Recorded backup clip
- ☐ Discovery questions printed

- ☐ Pilot proposal template with their facility name filled in
 - ☐ Calendar link for follow-up open
-

Tejas AI — the autonomous brain for cooling and energy. Built in India. Ready for the world.

Run the demo: `cd tejas-twin && ./run.sh` → <https://localhost:7878>